

How Much Does Limit Order Book Microstructure Add Beyond Quote Geometry?

Evidence from 70 days of cryptocurrency data at
retail-accessible sampling frequency

Jude Kriel Ramcharitar
Independent Researcher
`j.k.ramcharitar@gmail.com`

August 2026 · Version 4.0 — Working Paper

Abstract

A trader deciding where to post a limit order knows two things for free: how far from mid they intend to quote, and how long they are willing to wait. Whether limit order book (LOB) microstructure adds anything beyond those two choices is an empirical question, and one rarely posed against a baseline simple enough to make the answer interpretable.

We pose it directly. Using 269,330 LOB snapshots collected over a 70.5-day calendar window from Kraken across BTC/USD, ETH/USD and SOL/USD, we predict whether a limit order posted at a given distance from mid trades through within a fixed horizon, and we benchmark against a 16-cell lookup table containing only the mean historical fill rate for each combination of quote offset, horizon and side. That table is a demanding reference: quote offset spans a 0.46 range in fill rate and horizon a further 0.18.

Microstructure features beat it, modestly and consistently. Under a snapshot-grouped temporal split with a 990-second embargo, gradient-boosted trees over 63 microstructure features gain +0.048 and +0.060 AUC on BTC and ETH (Table 3 reports all figures to four decimals). Under a blocked split constructed so that training and test partitions span overlapping price regimes, the gains are +0.048, +0.049 and +0.027 on all three assets, with seed standard deviations at or below 0.0006 (Table 6). The mean across assets under blocking is +0.0415.

The split scheme matters, and revealed something we would otherwise have misreported. Under the temporal split, training and test partitions occupied almost disjoint price ranges: overlap of 0.129, 0.073 and 0.251. SOL/USD appeared to show a *negative* microstructure effect there, which reversed to +0.027 once the partitions spanned comparable price regimes. We report both schemes throughout, since neither substitutes for the other: the temporal split respects forecasting order, and the blocked split reduces temporal distribution shift between partitions.

We also compare model classes. Gradient-boosted trees outperform an LSTM on BTC by +0.040 and on ETH by +0.021, but the two are within 0.001 on SOL. An earlier version of this work claimed model class was decisive; with a correctly specified binary classification head, that claim does not hold.

Finally, we document five defects in our own pipeline, each of which produced plausible, publishable-looking numbers, and the specific checks that caught them. We consider this the most transferable part of the work.

Keywords: limit order book, fill probability, market microstructure, gradient boosting, baselines, reproducibility

1 Introduction

Consider a trader posting a limit buy order. Two quantities are entirely under their control and known before any model is consulted: the distance below mid at which they quote, and how long they will leave the order resting. Both bear strongly on whether the order fills. Quoting closer to mid fills more often; waiting longer fills more often.

The question this paper asks is what remains. Given those two choices, does the state of the order book — spreads, imbalances, depth, recent volatility — carry additional information about whether the order will fill?

The question is worth posing carefully because the obvious answer, that of course microstructure is informative, is not the same as a measured quantity. A large literature reports LOB models achieving substantial predictive accuracy [Zhang et al., 2019, Maglaras et al., 2022, Arroyo et al., 2024], but reported accuracy conflates what the model learned with what was mechanically available from quote placement. A model can post a respectable AUC while adding nothing a lookup table does not already provide.

1.1 What we find

Microstructure adds a modest, consistent amount. Under a blocked split that reduces temporal distribution shift between partitions, gradient-boosted trees over microstructure features gain +0.048, +0.049 and +0.027 AUC over the lookup reference on BTC/USD, ETH/USD and SOL/USD, mean +0.0415. Seed standard deviations are 0.0006 or below, so the ordering is stable across the three model seeds.

The finding is robust to split scheme on two of three assets and reverses favourably on the third. We report the mechanism in Section 6.

1.2 What we no longer claim

Version 3 of this work argued that model class was decisive: that trees recovered signal a deep survival model could not. That claim rested on a survival decoder being fed binary targets with a constant observed time, which pinned its training loss at 16.1181 and returned AUC of exactly 0.5000 on every arm. With a binary classification head the LSTM trains normally, and the separation between model classes is +0.040 on BTC, +0.021 on ETH and +0.001 on SOL. Trees are better here, but not categorically.

Versions 1 and 2 framed the work around execution latency, proposing that prediction could substitute for speed for traders operating at 200 ms–2 s delays. That framing does not survive: the labels supporting it were generated on a hardcoded 100 ms synthetic clock and, measured against real timestamps, the “2000 ms” horizon was a 511-second median window, an overstatement of 256×. The horizons we can actually resolve on this feed are 300 and 900 seconds, which is a different problem and is the one this paper addresses.

1.3 Contributions

1. A measurement of the marginal value of LOB microstructure over quote geometry, against an explicit trivial reference, on three assets under two split schemes (Sections 5, 6).

2. Evidence that a temporal split on trending data can invert the sign of a microstructure result, together with the diagnostic that detects it (Section 6).
3. A released dataset of 269,330 timestamped 10-level LOB snapshots with the full collection, labelling and evaluation pipeline (Section 3).
4. A catalogue of five pipeline defects that produced plausible results, and the checks that caught each (Section 7).

1.4 Scope

Established. In our Kraken sample, three cryptocurrency pairs from June to August 2026, microstructure features carry predictive content for our replay-defined fill outcome beyond quote offset, horizon and side, recoverable by gradient-boosted trees, at a magnitude of +0.027 to +0.049 AUC under a blocked split.

Not established: economic exploitability. AUC is not profit. We run no backtest, model no transaction costs, and make no claim that this edge survives the spread, maker fees, queue position or adverse selection.

Not established: realised fills. Labels are reconstructed from subsequent book states and do not model queue position. A quote is recorded as filled when top-of-book trades through its level, which is optimistic for an order arriving late in the queue.

Not established: calibration. AUC measures ranking, not whether a predicted 0.70 corresponds to a 70% fill rate.

Not established: forecasting performance. Our blocked-split results train on data from after the test period. They measure whether information is present, not whether it could have been traded in real time. The temporal results address that and are reported alongside.

Not established: generality. One venue, one 70-day window, three cryptocurrency pairs, one feed type, two horizons.

2 Related Work

2.1 Fill probability estimation

Lo et al. [2002] established econometric foundations for execution probability modelling. Maglaras et al. [2022] introduced a deep learning approach treating execution as a time-to-fill distribution. Arroyo et al. [2024] advanced this with a convolutional-transformer using survival analysis, which handles the right-censoring inherent in limit order data.

That work uses tick-level message data. Ours is snapshot data at median gaps of 11.5 to 48.0 seconds, three to four orders of magnitude coarser, so our results are not directly comparable and should not be read as contradicting theirs. What our setting permits, which theirs does not emphasise, is a clean separation between the mechanical contribution of quote placement and the informational contribution of book state.

2.2 Baselines in financial machine learning

Our methodological point — that a trivial reference is necessary to interpret a reported accuracy — is not novel in general, but is inconsistently applied in this literature. López de Prado [2018] discusses at length the ways financial backtests overstate performance, including through improper cross-validation on serially correlated data, and proposes purging and embargoing, which we adopt in Section 4.4.

2.3 Tabular deep learning

Systematic evaluations report gradient-boosted trees remaining competitive with neural networks on tabular data, particularly at moderate sample sizes with heterogeneous features [Grinsztajn et al., 2022, Shwartz-Ziv and Armon, 2022]. Our features are engineered per-snapshot scalars, which is tabular in structure. Our result is directionally consistent with that literature, though weaker than we previously claimed: trees win on two assets and tie on the third.

2.4 Market making theory

The framing that motivated earlier versions of this work derives from Avellaneda and Stoikov [2008], with latency treated as a structural risk factor by Cartea and Sánchez-Betancourt [2023] and Lehalle and Mounjid [2022]. We retain none of that framing here, because the horizons our data resolve are 300 and 900 seconds rather than the sub-second regime those models address. We note the connection only to be explicit that this paper does not extend it.

3 Data

3.1 Collection

LOB data were collected from Kraken via WebSocket API v2, subscribing to the `book` channel at depth 10 for BTC/USD, ETH/USD and SOL/USD. Updates are timestamped at receipt and stored with the full 10-level ladder in Parquet.

Collection ran from 15 June to 24 August 2026, a span of 70.5 days. The collector runs under a systemd unit with `Restart=always` and restarted automatically three times after WebSocket disconnections. An earlier configuration using a detached terminal session and a bounded reconnect loop failed silently on 20 July and stayed down for fourteen days.

Table 1: Collected data. Gaps are between consecutive book updates; the feed is event-driven, so sampling is irregular.

Asset	Snapshots	Tick	Median gap	Mean gap	Label rows
BTC/USD	111,002	0.10	11.5 s	54.9 s	1,760,728
ETH/USD	34,380	0.01	48.0 s	177.2 s	515,576
SOL/USD	123,948	0.01	26.2 s	49.1 s	1,971,128
Total	269,330				4,247,432

3.2 Labels

For each snapshot we simulate limit orders at offsets of 1, 2, 5 and 10 basis points from mid on each side, rounded to the venue tick so the quote is postable, and record whether top-of-book traded through that level within horizons of 300 and 900 seconds. A buy at price p is recorded filled if $\min(\text{best ask}) \leq p$ within the horizon; a sell if $\max(\text{best bid}) \geq p$.

Three choices here are deliberate and were wrong in earlier versions.

Horizons are resolved against real timestamps, not a fixed snapshot count. Earlier versions applied a hardcoded 100 ms interval inherited from a synthetic data generator (Section 7).

Offsets are in basis points, not ticks. Venue ticks are not economically comparable across these assets: BTC’s 0.10 tick on a $\sim 65,000$ price is 0.00015% of mid, while SOL’s 0.01 tick on ~ 150 is 0.0067%, roughly $44\times$ larger. Under tick offsets the 1–5 range spanned a 0.176 fill-rate gradient on SOL but only 0.011 on BTC, so “offset” meant different things per asset.

Unobservable rows are dropped, not labelled as non-fills. On a feed with a median gap of 11.5 to 48.0 seconds, some snapshots have no subsequent observation inside a horizon. Recording those as “did not fill” manufactures the label. Observability is 88.7% at worst (ETH at 300s) and 98.8% or above at 900s. This is why label counts are not snapshots $\times$ 16.

Table 2: Fill rate by horizon and quote offset. Both conditioning variables carry strong monotone structure, which is what makes the lookup reference demanding.

Asset	Horizon	1 bps	2 bps	5 bps	10 bps
BTC/USD	300 s	0.697	0.622	0.433	0.237
	900 s	0.815	0.766	0.626	0.440
ETH/USD	300 s	0.603	0.540	0.398	0.245
	900 s	0.740	0.696	0.582	0.432
SOL/USD	300 s	0.577	0.517	0.387	0.242
	900 s	0.753	0.712	0.612	0.466

3.3 Features

From each snapshot we construct 109 features: raw price and volume levels, normalised prices and volumes, mid-price, spread, spread in basis points, order book imbalance, depth imbalance, and rolling statistics and lagged differences.

Two subsets are used, both defined by keyword match on feature names. The *microstructure* subset holds the 63 features matching spread, imbalance, order book imbalance, depth, volume or volatility terms and not containing **price**. The *price* subset holds the 41 features whose names contain **price**, raw or normalised.

The two do not partition the feature set. Five features fall into neither: **return_1** and **return_lag_k** for $k \in \{1, 2, 3, 5\}$. These are lagged mid-price returns, so they are price-derived in substance, but the keyword rule matches the literal string **price** and does not capture them. The price subset is therefore *price levels*; returns enter only the full 109-feature arm.

We state this because $109 - 63 = 46$ counts non-microstructure columns rather than price columns, and the two differ by exactly these five returns. The microstructure subset is unaffected: returns match no microstructure keyword either, so those 63 features contain no price information in level or change form. Headline results use that subset, for reasons given in Section 6.

4 Method

4.1 The reference

Our benchmark is a lookup table: the mean fill rate for each (offset, horizon, side) cell, 16 cells, estimated on training data and applied unchanged to test data. It has no learned parameters and no access to the order book. Table 2 shows what it captures.

4.2 Models

Gradient-boosted trees (XGBoost, 400 estimators, depth 6, learning rate 0.05, subsample 0.8) on the current snapshot’s features plus the three conditioning variables. No temporal window.

Binary LSTM: two layers, hidden dimension 64, over a 50-snapshot window, with the conditioning vector concatenated before a single logit head trained with binary cross-entropy. The comparison is deliberately unfavourable to the trees, which see one snapshot where the LSTM sees fifty.

4.3 Sampling

Model fits do not use all 4,247,432 label rows. For each asset we draw a uniform random sample without replacement of 200,000 rows, using a fixed seed (0), and restore temporal order before splitting. The cap bounds LSTM training time, and the same sampled set is used by every arm so the comparison is on identical rows.

Sampling occurs before the split, so partition proportions are computed on the sampled rows. This is 11.4% of available BTC/USD labels, 38.8% of ETH/USD and 10.1% of SOL/USD, so the effective sample is more nearly balanced across assets than Table 1 suggests. Because sampling is uniform over rows rather than stratified by snapshot, it also thins the number of label rows per snapshot; the grouped split (below) still operates on whole snapshots.

4.4 Splitting and embargo

Each snapshot generates up to 16 label rows sharing identical features, so a row-wise split permits a snapshot to appear in both training and test data. All splits are therefore *snapshot-grouped*.

Labels look up to 900 seconds forward. Every partition boundary carries a 990-second embargo, following the purging logic of López de Prado [2018], so a training label’s forward window cannot reach into test data.

We report two schemes. The **temporal** split takes the first 70% of snapshots for training, the next 15% for validation and the last 15% for test, which respects forecasting order. The **blocked** split cuts the timeline into 150 contiguous blocks assigned in a repeating pattern, so all partitions span the full price history. Section 6 explains why both are necessary.

4.5 Metric

We report fill AUC: area under the ROC curve for the binary filled/not-filled task, computed exactly via the Mann-Whitney rank-sum identity with ties credited 0.5. It measures discrimination. It does not assess calibration, so these results do not establish that predicted probabilities are usable as probabilities.

5 Results

5.1 Main comparison

Seed standard deviations, shown beneath each mean, span 0.0006–0.0035 for GBT arms and 0.0021–0.0074 for LSTM arms, so the differences discussed below exceed run-to-run variation. They quantify model-initialisation variance only, not variance over data partitions.

Two observations, and one anomaly deferred to Section 6.

Table 3: Fill AUC under the temporal split, three seeds. Each cell gives the mean, the seed standard deviation, and the gap over the lookup reference for that asset. GBT sees one snapshot; LSTM sees fifty.

Asset	Lookup	GBT micro (63)	GBT all (109)	LSTM micro (63)	LSTM all (109)
BTC/USD	0.6828	0.7307	0.7262	0.6857	0.6903
		± 0.0006	± 0.0012	± 0.0021	± 0.0043
		+0.0479	+0.0434	+0.0029	+0.0075
ETH/USD	0.6597	0.7193	0.6908	0.6983	0.6669
		± 0.0013	± 0.0010	± 0.0046	± 0.0036
		+0.0596	+0.0311	+0.0386	+0.0072
SOL/USD	0.6673	0.6568	0.7056	0.7043	0.7043
		± 0.0035	± 0.0010	± 0.0046	± 0.0074
		-0.0105	+0.0383	+0.0370	+0.0370

Microstructure features add measurable signal on BTC and ETH: +0.048 and +0.060 for the tree model on the 63-feature subset. SOL shows -0.011, the only negative entry in the table, which we resolve below.

Adding the 46 non-microstructure columns (41 price levels and five lagged returns) *hurts* on BTC and ETH, with GBT-all below GBT-micro by 0.005 and 0.028. On SOL they appear to help substantially. This inconsistency is the same anomaly.

5.2 Model class

Comparing the better arm of each model class:

Table 4: Best gap over lookup by model class, temporal split.

Asset	Best GBT	Best LSTM	Difference
BTC/USD	+0.0479	+0.0075	+0.0404
ETH/USD	+0.0596	+0.0386	+0.0210
SOL/USD	+0.0383	+0.0370	+0.0013

Trees are ahead on all three, decisively on BTC, marginally on SOL. Version 3 of this work reported the LSTM at exactly 0.5000 on every arm and concluded model class was decisive; that was a defective survival head, not a result (Section 7). The corrected comparison is narrower: trees are better on this data, and on one asset the difference is 0.0013.

One arm warrants caution. LSTM-all on ETH moved its training loss by only 0.0078, against 0.1593 for LSTM-micro on the same data. That arm is likely undertrained rather than genuinely uninformative, and we do not draw a feature-count conclusion from it.

6 The Split Scheme Inverts a Result

6.1 Diagnosis

SOL’s negative microstructure gap prompted a check on how the temporal split partitions price levels. The result applies to all three assets.

Table 5: Overlap between training and test price ranges, by split scheme. Low overlap means absolute price nearly identifies the partition.

Asset	Temporal	Blocked
BTC/USD	0.129	0.815
ETH/USD	0.073	0.855
SOL/USD	0.251	0.745

Over a trending 70-day window, a temporal split leaves training and test partitions occupying nearly disjoint price ranges. ETH trains on \$1510–1944 and tests on \$1869–2545, an overlap of 0.073. A tree splitting on absolute price can identify the partition almost perfectly, which makes any gain attributable to price-derived features difficult to interpret.

The blocked split cuts the timeline into 150 contiguous blocks assigned in a repeating pattern with the same 990-second embargo at every boundary, raising overlap to 0.745–0.855.

6.2 Effect

Table 6: Gap over lookup by feature group and split scheme, three seeds. For the blocked microstructure arm the underlying AUCs are 0.7519, 0.7141 and 0.6955 against lookup references of 0.7040, 0.6650 and 0.6681, with seed standard deviations of 0.0004, 0.0006 and 0.0003.

Asset	Microstructure (63)		Price (41)		All (109)	
	Temporal	Blocked	Temporal	Blocked	Temporal	Blocked
BTC/USD	+0.0479	+0.0479	+0.0384	+0.0383	+0.0434	+0.0456
ETH/USD	+0.0596	+0.0491	−0.0001	+0.0370	+0.0312	+0.0513
SOL/USD	−0.0105	+0.0274	+0.0312	+0.0115	+0.0383	+0.0295

Three findings.

The microstructure result is robust where it was already positive. BTC is identical to four decimals (+0.0479 under both); ETH holds within 0.011.

SOL’s negative result is sensitive to the split scheme. It moves from −0.0105 under the forecasting-ordered split to +0.0274 once partitions span comparable price regimes. We report this as sensitivity to temporal distribution shift rather than as a demonstrated price-level mechanism: blocking changes several things at once, including price-range overlap, regime composition, temporal distance between partitions, and whether training data may post-date test data. We have not isolated which of these drives the reversal. With blocking, all three assets show a positive microstructure gap, ranging +0.027 to +0.049 with a mean of +0.0415.

Price features behave inconsistently and we do not rely on them. On ETH they move from −0.000 to +0.037 under blocking; on SOL from +0.031 to +0.012. We had expected blocking to reduce price-feature performance uniformly by removing a time proxy. It did not, and we have no account of the ETH reversal. Headline figures therefore use the microstructure subset, which excludes price-derived columns entirely and behaves consistently under both schemes.

6.3 Which scheme to believe

Neither replaces the other. The temporal split respects forecasting order and is the appropriate evaluation for a strategy question; the blocked split reduces temporal distribution shift between partitions, at the cost noted in Section 1.4. That the microstructure result holds under both on BTC and ETH is why we report it as a finding rather than a split-specific number.

7 Five Defects

Each of the following produced a plausible result. None was detectable by inspecting output. We list them with the check that caught each, because that is the transferable content.

1. Quote offset never reached the model. The conditioning vector carried only the latency window. For a fixed snapshot, label rows spanning four offsets presented identical model input with different targets, so the loss minimiser was to predict the mean and ignore the input. Symptom: validation loss moving 0.003 over 40 epochs. Effect: C-index 0.4749, rising to 0.6266 once corrected. *Caught by* a diagnostic asking whether a lookup table on the conditioning variables alone could beat the model. It could.

2. A synthetic 100 ms clock applied to real data. The label builder converted a latency window to a fixed snapshot count using an interval hardcoded for the synthetic generator. On real data, the “200 ms” label was a 33.5-second median horizon and “2000 ms” was 511.1 seconds, overstatements of $168\times$ and $256\times$. *Caught by* an arithmetic fingerprint. That version used a $4 \times 4 \times 2 = 32$ label grid (four offsets, four latency windows, two sides), against the $4 \times 2 \times 2 = 16$ grid used in this paper. Reported label counts fell short of snapshots $\times 32$ by exactly 672 rows on every asset, which is 21×32 , and 21 is the tail a loop bounded at $n - 2000/100 - 1$ leaves behind. A horizon resolved against real timestamps would have left a different remainder per asset, since their gap distributions differ.

3. Microsecond timestamps read as nanoseconds. The replacement label builder converted timestamps with `astype("int64") / 1e9`, but Parquet returns `datetime64[us]`. Every interval came out $1000\times$ too small, so a “30-second” window contained 920 snapshots and fill rates reached 0.97 with no offset gradient. *Caught by* a printed diagnostic showing a median gap of 0.0 s against a known 11.5 s. The current builder raises if the median gap is below 1 second.

4. A survival head fed binary targets. When labels became binary, they were passed to a time-to-fill decoder with observed time set to a constant. Training loss sat at 16.1181 from the first epoch and test AUC was exactly 0.5000 on all six arms. *Caught by* a positive control: the same architecture with a binary head reaches AUC 0.9861 on synthetic data with a planted signal, so the failure was the head, not the data.

5. Split proportions computed on time rather than count. Cutting the time range at 70/85% gave partitions of 44/31/25% because snapshot density varies across the window. *Caught by* asserting the realised proportions.

7.1 What generalises

The common feature is that every defect produced output within the range a reader would accept. C-index of 0.5526 looks like a weak model. AUC of 0.8264 looks like a strong one. Neither was real.

Three practices follow. Run the trivial reference before the model, since four of the five defects would have been visible as a model failing to beat a lookup table. Write a positive control on

planted signal, since defect 4 is undetectable any other way. Assert the properties of your splits rather than assuming them.

We also note the cost of misattributing a failure. After defect 1, we read a flat learning curve as a sample-size problem, projected that 300,000 snapshots would be required, and collected for six additional weeks. The projection was extrapolated from noise generated by the defect.

8 Limitations

No economic evaluation. AUC is not profit, and we do not model transaction costs, queue position or adverse selection.

Simulated fills. Labels come from replaying book states. A quote is filled when top-of-book trades through its level, which is optimistic for a late queue arrival.

Two horizons. 300 and 900 seconds, chosen because observability at 30 and 60 seconds fell to 41–54% on ETH and SOL.

One venue, one window, three assets.

One deep architecture. An LSTM at one hyperparameter setting. The ETH LSTM-all arm is likely undertrained.

Partition variance is not quantified. Seeds vary model initialisation, not the data partition. We evaluate two partition schemes (Section 6) rather than a distribution over partitions, so the reported standard deviations understate total uncertainty.

The price-feature behaviour is unexplained. We report it and exclude those features from headline results rather than offering an account we cannot support.

9 Conclusion

Given a quote offset and a holding horizon — both known to the trader for free — limit order book microstructure adds roughly 0.04 AUC to fill prediction on this data. Under a blocked split the gain is +0.048, +0.049 and +0.027 across three cryptocurrency pairs, with seed variation an order of magnitude smaller than the effect. Under a forecasting-ordered temporal split it holds on two of three assets, and the third is sensitive to temporal distribution shift rather than showing absent signal.

That is a real effect and a small one. It is smaller than the numbers we reported in three previous versions of this work, each of which was inflated by a defect in our own pipeline. We consider the catalogue in Section 7 at least as useful as the measurement, because the defects were not exotic and the outputs they produced were not obviously wrong.

The question of whether a 0.04 AUC edge over a lookup table is worth anything after costs is the natural next one, and this paper does not answer it.

Data and Code Availability

The collection pipeline, label builder, models, verification code and every diagnostic reported here are available at <https://github.com/azure-cyber/latency-compensating-mm>. The verification harness (`verify_v3.py`) runs the positive control and split assertions described in Section 7 and exits non-zero on failure.

Disclosure of AI Assistance

The author used Claude (Anthropic) extensively throughout this project: for implementation of the data collection pipeline, feature engineering, models and evaluation code; for diagnostic analysis; and for drafting and revision of this manuscript. Several of the defects catalogued in Section 7 originated in AI-generated code, and several were identified through AI-assisted diagnosis. The author designed the study, directed the analysis, made all substantive research decisions, verified the results, and is solely responsible for the content and any errors.

Acknowledgements

This work was conducted independently, without institutional affiliation, funding or supervision. The author declares no conflicts of interest and holds no positions in any asset discussed.

References

- Álvaro Arroyo, Álvaro Cartea, Fernando Moreno-Pino, and Stefan Zohren. Deep attentive survival analysis in limit order books: Estimating fill probabilities with convolutional-transformers. *Quantitative Finance*, 24(1):35–57, 2024.
- Marco Avellaneda and Sasha Stoikov. High-frequency trading in a limit order book. *Quantitative Finance*, 8(3):217–224, 2008.
- Álvaro Cartea and Leandro Sánchez-Betancourt. Optimal execution with stochastic delay. *Finance and Stochastics*, 27:1–47, 2023.
- Léo Grinsztajn, Edouard Oyallon, and Gaël Varoquaux. Why do tree-based models still outperform deep learning on typical tabular data? In *Advances in Neural Information Processing Systems (Datasets and Benchmarks Track)*, 2022.
- Charles-Albert Lehalle and Othmane Mounjid. Limit order strategic placement with adverse selection risk and the role of latency. *Market Microstructure and Liquidity*, 2022. arXiv:1610.00261.
- Andrew W. Lo, A. Craig MacKinlay, and June Zhang. Econometric models of limit-order executions. *Journal of Financial Economics*, 65(1):31–71, 2002.
- Marcos López de Prado. *Advances in Financial Machine Learning*. Wiley, 2018.
- Costis Maglaras, Ciamac C. Moallemi, and Muye Wang. A deep learning approach to estimating fill probabilities in a limit order book. *Quantitative Finance*, 22(11):1989–2003, 2022.
- Ravid Shwartz-Ziv and Amitai Armon. Tabular data: Deep learning is not all you need. *Information Fusion*, 81:84–90, 2022.
- Zihao Zhang, Stefan Zohren, and Stephen Roberts. DeepLOB: Deep convolutional neural networks for limit order books. *IEEE Transactions on Signal Processing*, 67(11):3001–3012, 2019.